

create index idx_customer_id ON dim_customers(customer_id).
clustered vs Non-clustered


- points back to the original rows.
- like the back of the book index

Physical reorders the rows of the table itself to match the index order

Databricks Toolkit

- Instead of index point to individual rows databricks works at a file level.
- Skips files that don't contain relevant data.

Partition by → Splits data into folders by column value
Optimize → Compacts small files



Z order → group related rows into the same files

So

Subqueries

Select statement nested inside another select

A

Select

Customer ID, Customer name, Customer age

From (select customer_id

From dim_customer

where age > 25)

}

Inner query runs first